



# Keeping your GPU fed without getting bitten

Tobias Hector  
May 2016


# Introduction

- You have delicious draw calls
  - Yummy!




# Introduction

- You have delicious draw calls
  - Yummy!
- Your GPU wants to eat them
  - It's really hungry



# Introduction

- You have delicious draw calls
  - Yummy!
- Your GPU wants to eat them
  - It's really hungry
- Keep it fed at all times
  - So it keeps making pixels



# Introduction

- You have delicious draw calls
  - Yummy!
- Your GPU wants to eat them
  - It's really hungry
- Keep it fed at all times
  - So it keeps making pixels
- Don't want it biting your hand
  - Look at those teeth!



# Keeping it fed

- GPU needs a constant supply of food
  - It doesn't want to wait
- Certain foods are tough to digest
  - Provide multiple operations to hide stalls
- Draw calls provide a variety of nutrition
  - Vertex work, raster work, tessellation, vitamins A-K, etc.

# Keeping it fed

System			
CPU	0		1
GPU		0	1

# Keeping it fed

System			
CPU	0	1	2
GPU		0	1



# Keeping it fed

GPU				
Vertex	0		1	
Fragment		0		1

# Keeping it fed

GPU				
Vertex	0	1	2	
Fragment		0	1	2

# Not getting bitten

- GPU eating from lots of different plates
  - Don't touch anything it's using!
- It doesn't want a mouthful of beef choc chip ice cream
  - Don't change data whilst it's accessing a resource
- Hey I'm eating that!
  - Don't delete resources whilst the GPU is still using them






Tear Point #1 --->

<sup>其</sup>  
Tear Point #2 --->















# On to the serious bits...

# Terminology

- **Operation**
  - Anything that can be executed
    - Includes synchronization and memory barriers
- **Execution Dependency**
  - Operations waiting on other operations
  - All synchronization expresses these
- **Memory Barrier**
  - Flush/invalidate caches
  - Determination of access and visibility
- **Memory Dependency**
  - Execution dependency involving a Memory Barrier

Note: Memory barrier does not mean quite the same thing as GL's memory barrier, though there is some relation.

# Synchronization Types

- 3 types of explicit synchronization in Vulkan
  - Pipeline Barriers, Events and Subpass Dependencies
    - Within a queue
    - Explicit memory dependencies
  - Semaphores
    - Between Queues
  - Fences
    - Whole queue operations to CPU

OpenGL has just two, very coarse synchronization primitives: memory barriers and fences. They are loosely similar to the equivalently named concepts in Vulkan

# Pipeline Barriers

- Pipeline Barriers
  - Precise set of pipeline stages
  - Memory Barriers to execute
  - Single point in time

Executing a pipeline barrier is roughly equivalent to a `glMemoryBarrier` call, though with much more control.

```
void vkCmdPipelineBarrier(  
    VkCommandBuffer          commandBuffer,  
    VkPipelineStageFlags     srcStageMask,  
    VkPipelineStageFlags     dstStageMask,  
    VkDependencyFlags        dependencyFlags,  
    uint32_t                 memoryBarrierCount,  
    const VkMemoryBarrier*   pMemoryBarriers,  
    uint32_t                 bufferMemoryBarrierCount,  
    const VkBufferMemoryBarrier* pBufferMemoryBarriers,  
    uint32_t                 imageMemoryBarrierCount,  
    const VkImageMemoryBarrier* pImageMemoryBarriers);
```

# Events

- Events
  - Same info as Pipeline Barriers
  - ...but operate over a range

```
void vkCmdSetEvent(  
    VkCommandBuffer      commandBuffer,  
    VkEvent               event,  
    VkPipelineStageFlags stageMask);  
  
void vkCmdResetEvent(  
    VkCommandBuffer      commandBuffer,  
    VkEvent               event,  
    VkPipelineStageFlags stageMask);  
  
void vkCmdWaitEvents(  
    VkCommandBuffer      commandBuffer,  
    uint32_t              eventCount,  
    const VkEvent*        pEvents,  
    VkPipelineStageFlags srcStageMask,  
    VkPipelineStageFlags dstStageMask,  
    uint32_t              memoryBarrierCount,  
    const VkMemoryBarrier* pMemoryBarriers,  
    uint32_t              bufferMemoryBarrierCount,  
    const VkBufferMemoryBarrier* pBufferMemoryBarriers,  
    uint32_t              imageMemoryBarrierCount,  
    const VkImageMemoryBarrier* pImageMemoryBarriers);
```

# Events

- **Events**
  - Same info as Pipeline Barriers
  - ...but operate over a range
- **CPU interaction**
  - No explicit CPU wait
  - No Memory Barriers

```
VkResult vkSetEvent(  
    VkDevice      device,  
    VkEvent       event);
```

```
VkResult vkResetEvent(  
    VkDevice      device,  
    VkEvent       event);
```

```
VkResult vkGetEventStatus(  
    VkDevice      device,  
    VkEvent       event);
```

# Events

- **Events**

- Same info as Pipeline Barriers
- ...but operate over a range

- **CPU interaction**

- No explicit CPU wait
- No Memory Barriers

- **Warning!**

- OS may apply a timeout
- Set events soon after submission
- Could you just defer submission?

```
VkResult vkSetEvent(  
    VkDevice device,  
    VkEvent event);
```

```
VkResult vkResetEvent(  
    VkDevice device,  
    VkEvent event);
```

```
VkResult vkGetEventStatus(  
    VkDevice device,  
    VkEvent event);
```



# Pipeline Barriers vs Events

- **Use pipeline barriers for point synchronization**
  - Dependant operation immediately precedes operation that depends on it
  - May be more optimal than set/wait event pair
- **Use events if other work possible between two operations**
  - Set immediately after the dependant operation
  - Wait immediately before the operation that depends on it
- **Use events for CPU/GPU synchronization**
  - Memory accesses between processors
  - Late latching of data to reduce latency



# Memory Barrier Types

- **Global Memory Barrier**
  - All memory-backed resources
- **Buffer Barrier**
  - For a single buffer range
- **Image Barrier**
  - For a single image subresource range

OpenGL's memory barriers imply execution dependencies, which Vulkan memory barriers do not - execution barriers are provided by a pipeline barrier, event or subpass dependency.

# Global Memory Barriers

- Global Memory Barriers
  - All memory used by **accessed stages**
  - Effectively flushes entire caches
- Use when many resources transition
  - Cheaper than one-by-one
  - Don't transition unnecessarily!
- User must define prior access
  - Driver not tracking for you

```
typedef struct VkMemoryBarrier {  
    VkStructureType      sType;  
    const void*          pNext;  
    VkAccessFlags         srcAccessMask;  
    VkAccessFlags         dstAccessMask;  
} VkMemoryBarrier;
```

# Buffer Barriers

- Buffer Barriers
  - A single **buffer range**
  - Defines **access stages**
  - Defines **queue ownership**
- User must define prior access
  - Driver not tracking for you

```
typedef struct VkBufferMemoryBarrier {  
    VkStructureType      sType;  
    const void*          pNext;  
    VkAccessFlags         srcAccessMask;  
    VkAccessFlags         dstAccessMask;  
    uint32_t              srcQueueFamilyIndex;  
    uint32_t              dstQueueFamilyIndex;  
    VkBuffer              buffer;  
    VkDeviceSize          offset;  
    VkDeviceSize          size;  
} VkBufferMemoryBarrier;
```

# Image Barriers

- Image Barriers
  - A single **image subresource range**
  - Defines **access stages**
  - Defines **queue ownership**
  - Defines **image layout**
- User must define prior access
  - Driver not tracking for you
  - For images, this includes prior layout
- Appropriate layouts allow compression
  - GPU may use image compression
  - Saves bandwidth
  - Use GENERAL instead of switching frequently

```
typedef struct VkImageMemoryBarrier {  
    VkStructureType      sType;  
    const void*          pNext;  
    VkAccessFlags         srcAccessMask;  
    VkAccessFlags         dstAccessMask;  
    VkImageLayout         oldLayout;  
    VkImageLayout         newLayout;  
    uint32_t             srcQueueFamilyIndex;  
    uint32_t             dstQueueFamilyIndex;  
    VkImage               image;  
    VkImageSubresourceRange subresourceRange;  
} VkImageMemoryBarrier;
```

# Subpass Dependencies

- Subpass dependencies
  - Similar info to Pipeline Barriers
  - Explicitly between **two subpasses**
- Memory barriers
  - Implicit for attachments
  - **Explicit for other resources**
- Pixel local dependencies
  - Same fragment/sample location
  - Cheap for most implementations
  - Use region dependency flag:
    - VK\_DEPENDENCY\_BY\_REGION\_BIT

```
typedef struct VkSubpassDependency {  
    uint32_t                srcSubpass;  
    uint32_t                dstSubpass;  
    VkPipelineStageFlags    srcStageMask;  
    VkPipelineStageFlags    dstStageMask;  
    VkAccessFlags           srcAccessMask;  
    VkAccessFlags           dstAccessMask;  
    VkDependencyFlags       dependencyFlags;  
} VkSubpassDependency;
```

# Subpass Dependencies

- Subpass self-dependencies
  - Subpasses can wait on themselves
  - A pipeline barrier in the subpass
- Forward progress only
  - Can't wait on later stages
  - Must wait on earlier or same stage
- Pixel local only between fragments
  - Must use flag:
    - VK\_DEPENDENCY\_BY\_REGION\_BIT

```
typedef struct VkSubpassDependency {
    uint32_t          srcSubpass;
    uint32_t          dstSubpass;
    VkPipelineStageFlags srcStageMask;
    VkPipelineStageFlags dstStageMask;
    VkAccessFlags      srcAccessMask;
    VkAccessFlags      dstAccessMask;
    VkDependencyFlags   dependencyFlags;
} VkSubpassDependency;

void vkCmdPipelineBarrier(
    VkCommandBuffer          commandBuffer,
    VkPipelineStageFlags     srcStageMask,
    VkPipelineStageFlags     dstStageMask,
    VkDependencyFlags        dependencyFlags,
    uint32_t                 memoryBarrierCount,
    const VkMemoryBarrier*   pMemoryBarriers,
    uint32_t                 bufferMemoryBarrierCount,
    const VkBufferMemoryBarrier* pBufferMemoryBarriers,
    uint32_t                 imageMemoryBarrierCount,
    const VkImageMemoryBarrier* pImageMemoryBarriers);
```

# Subpass Dependencies

- Subpass external dependencies
  - Wait on 'external' operations
  - vkCmdWaitEvent in the subpass
  - Events set outside the render pass

```
typedef struct VkSubpassDependency {
    uint32_t                srcSubpass;
    uint32_t                dstSubpass;
    VkPipelineStageFlags    srcStageMask;
    VkPipelineStageFlags    dstStageMask;
    VkAccessFlags           srcAccessMask;
    VkAccessFlags           dstAccessMask;
    VkDependencyFlags       dependencyFlags;
} VkSubpassDependency;

void vkCmdWaitEvents(
    VkCommandBuffer          commandBuffer,
    uint32_t                eventCount,
    const VkEvent*           pEvents,
    VkPipelineStageFlags     srcStageMask,
    VkPipelineStageFlags     dstStageMask,
    uint32_t                memoryBarrierCount,
    const VkMemoryBarrier*   pMemoryBarriers,
    uint32_t                bufferMemoryBarrierCount,
    const VkBufferMemoryBarrier* pBufferMemoryBarriers,
    uint32_t                imageMemoryBarrierCount,
    const VkImageMemoryBarrier* pImageMemoryBarriers);
```

# Example - Texture Upload

```
// Transition the buffer from host write to transfer read
bufferBarrier.srcAccessMask = VK_ACCESS_HOST_WRITE_BIT;
bufferBarrier.dstAccessMask = VK_ACCESS_TRANSFER_READ_BIT;
// Transition the image to transfer destination
imageBarrier.srcAccessMask = 0;
imageBarrier.dstAccessMask = VK_ACCESS_TRANSFER_WRITE_BIT;
imageBarrier.oldLayout = VK_IMAGE_LAYOUT_UNDEFINED;
imageBarrier.newLayout = VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL;

vkCmdPipelineBarrier(commandBuffer, VK_PIPELINE_STAGE_HOST_BIT, VK_PIPELINE_STAGE_TRANSFER_BIT, &bufferBarrier,
&imageBarrier);

vkCmdCopyBufferToImage(commandBuffer, srcBuffer, image, VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL, 1, &copy);

// Transition the image from transfer destination to shader read
imageBarrier.srcAccessMask = VK_ACCESS_TRANSFER_WRITE_BIT;
imageBarrier.dstAccessMask = VK_ACCESS_SHADER_READ_BIT;
imageBarrier.oldLayout = VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL;
imageBarrier.newLayout = VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL;

vkCmdPipelineBarrier(commandBuffer, VK_PIPELINE_STAGE_TRANSFER_BIT, VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,
&imageBarrier);
```



# Example - Compute to Draw Indirect

```
// Add a subpass dependency to express the wait on an external event
externalDependency.srcSubpass = VK_SUBPASS_EXTERNAL;
externalDependency.srcStageMask = VK_PIPELINE_STAGE_COMPUTE_SHADER_BIT;
externalDependency.dstStageMask = VK_PIPELINE_STAGE_DRAW_INDIRECT_BIT;
externalDependency.srcAccessMask = VK_ACCESS_SHADER_WRITE_BIT;
externalDependency.dstAccessMask = VK_ACCESS_INDIRECT_COMMAND_READ_BIT;
```

```
// Dispatch a compute shader that generates indirect command structures
vkCmdDispatch(...);
// Set an event that can be later waited on (same source stage).
vkCmdSetEvent(commandBuffer, event, VK_PIPELINE_STAGE_COMPUTE_SHADER_BIT);
```

```
vkCmdBeginRenderPass(...);

//Transition the buffer from shader write to indirect command
bufferBarrier.srcAccessMask = VK_ACCESS_SHADER_WRITE_BIT;
bufferBarrier.dstAccessMask = VK_ACCESS_INDIRECT_COMMAND_READ_BIT;
bufferBarrier.buffer = indirectBuffer;
vkCmdWaitEvent(commandBuffer, event, VK_PIPELINE_STAGE_COMPUTE_SHADER_BIT, VK_PIPELINE_STAGE_DRAW_INDIRECT_BIT,
&bufferBarrier);

vkCmdDrawIndirect(commandBuffer, indirectBuffer, ...);
```

# Semaphores

- **Semaphores**
  - Used to synchronize queues
  - Not necessary for single-queue
- **Fairly coarse grain**
  - Per submission batch
    - E.g. a set of command buffers
  - Multiple per submit command
- **Implicit memory guarantees**
  - Effects visible to future operations on the same device
    - Not guaranteed visible to host

```
typedef struct VkSubmitInfo {  
    VkStructureType      sType;  
    const void*          pNext;  
    uint32_t             waitSemaphoreCount;  
    const VkSemaphore*    pWaitSemaphores;  
    const VkPipelineStageFlags* pWaitDstStageMask;  
    uint32_t             commandBufferCount;  
    const VkCommandBuffer* pCommandBuffers;  
    uint32_t             signalSemaphoreCount;  
    const VkSemaphore*    pSignalSemaphores;  
} VkSubmitInfo;
```

# Example - Acquire and Present

```
// Acquire an image. Pass in a semaphore to be signalled
```

```
vkAcquireNextImageKHR(device, swapchain, UINT64_MAX, acquireSemaphore, VK_NULL_HANDLE, &imageIndex);
```

```
// Submit command buffers
```

```
submitInfo.waitSemaphoreCount = 1;
```

```
submitInfo.pWaitSemaphores = &acquireSemaphore;
```

```
submitInfo.commandBufferCount = 1;
```

```
submitInfo.pCommandBuffers = &commandBuffer;
```

```
submitInfo.signalSemaphoreCount = 1;
```

```
submitInfo.pWaitSemaphores = &graphicsSemaphore;
```

```
vkQueueSubmit(graphicsQueue, 1, &submitInfo, fence);
```

```
// Present images to the display
```

```
presentInfo.waitSemaphoreCount = 1;
```

```
presentInfo.pWaitSemaphores = &graphicsSemaphore;
```

```
presentInfo.swapchainCount = 1;
```

```
presentInfo.pSwapchains = &swapchain;
```

```
presentInfo.pImageIndices = &imageIndex;
```

```
vkQueuePresent(presentQueue, &presentInfo);
```

# Example - Acquire and Present (same queue)

```
// Acquire an image. Pass in a semaphore to be signalled
```

```
vkAcquireNextImageKHR(device, swapchain, UINT64_MAX, acquireSemaphore, VK_NULL_HANDLE, &imageIndex);
```

```
// Submit command buffers
```

```
submitInfo.waitSemaphoreCount = 1;
```

```
submitInfo.pWaitSemaphores = &acquireSemaphore;
```

```
submitInfo.commandBufferCount = 1;
```

```
submitInfo.pCommandBuffers = &commandBuffer;
```

```
submitInfo.signalSemaphoreCount = 0;
```

```
vkQueueSubmit(universalQueue, 1, &submitInfo, fence);
```

```
// Present images to the display
```

```
presentInfo.waitSemaphoreCount = 0;
```

```
presentInfo.swapchainCount = 1;
```

```
presentInfo.pSwapchains = &swapchain;
```

```
presentInfo.pImageIndices = &imageIndex;
```

```
vkQueuePresent(universalQueue, &presentInfo);
```

# Fences

- Fences
  - Used to synchronize queue to CPU
- Very coarse grain
  - Per queue submit command
- Implicit memory guarantees
  - Effects visible to future operations on the same device
    - Not guaranteed visible to host

GL's fences are like a combination of a semaphore and a fence in Vulkan - they can synchronize GPU and CPU in multiple ways at a coarse granularity.

```
VkResult vkQueueSubmit(  
    VkQueue          queue,  
    uint32_t         submitCount,  
    const VkSubmitInfo* pSubmits,  
    VkFence          fence);
```

```
VkResult vkResetFences(  
    VkDevice          device,  
    uint32_t         fenceCount,  
    const VkFence*    pFences);
```

```
VkResult vkGetFenceStatus(  
    VkDevice          device,  
    VkFence          fence);
```

```
VkResult vkWaitForFences(  
    VkDevice          device,  
    uint32_t         fenceCount,  
    const VkFence*    pFences,  
    VkBool32          waitAll,  
    uint64_t          timeout);
```

# Example - Multi-buffering

```
// Have enough resources and fences to have one per in-flight-frame, usually the swapchain image count
VkBuffer buffers[swapchainImageCount];
VkFence fence[swapchainImageCount];

// Can use the index from the presentation engine - 1:1 mapping between swapchain images and resources
vkAcquireNextImageKHR(device, swapchain, UINT64_MAX, semaphore, VK_NULL_HANDLE, &nextIndex);

// Make absolutely sure that the work has completed
vkWaitForFences(device, 1, &fence[nextIndex], true, UINT64_MAX);

// Reset the fences we waited on, so they can be re-used
vkResetFences(device, 1, &fence[nextIndex]);

// Change the data in your per-frame resources (with appropriate events/barriers!)
...

// Submit any work to the queue, with those fences being re-used for the next time around
vkQueueSubmit(graphicsQueue, 1, &sSubmitInfo, fence[nextIndex]);
```

# Wait Idle

- Ensures execution completes
  - VERY heavy-weight
- **vkQueueWaitIdle**
  - Wait for queue operations to finish
  - Equivalent to waiting on a fence
- **vkDeviceWaitIdle**
  - Waits for device operations to finish
  - Includes vkQueueWaitIdle for queues

These are a lot like glFinish, and should be treated similarly - use them VERY SPARINGLY.

```

VkResult vkQueueSubmit(
    VkQueue          queue,
    uint32_t         submitCount,
    const VkSubmitInfo* pSubmits,
    VkFence          fence);

VkResult vkResetFences(
    VkDevice          device,
    uint32_t          fenceCount,
    const VkFence*    pFences);

VkResult vkGetFenceStatus(
    VkDevice          device,
    VkFence           fence);

VkResult vkWaitForFences(
    VkDevice          device,
    uint32_t          fenceCount,
    const VkFence*    pFences,
    VkBool32          waitAll,
    uint64_t          timeout);

```

# Wait Idle

- **Useful primarily at teardown**
  - Use it to quickly ensure all work is done
- **Favour other synchronization at all other times**
  - Extremely heavyweight, will cause serialization!



# Programmer Guidelines

- **Specify EXACTLY the right amount of synchronization**
  - Too much and you risk starving your GPU
  - Miss any and your GPU will bite you
- **Use the validation layers to help!**
  - Won't catch everything yet, but improving over time
- **Pay particular attention to the pipeline stages**
  - Fiddly but become intuitive as you use them
- **Consider Image Layouts**
  - If your GPU can save bandwidth it will
- **Different behaviour depending on implementation**
  - Test/Tune on every platform you can find!

# Keep your GPU fed without getting bitten!

## Questions?